

CSCI235 – Database Systems

Introduction to Indexing

sjapit@uow.edu.au

1 October 2021



Acknowledgements

The following presentation were adapted from the lecture slides of:

CSCI235 – Database Systems, 06Introduction to Indexing

By Dr Janusz R. Getta, University of Wollongong, Australia

Spring, 2018

Outline

- Index? What is it?
- Index versus indexed file organization
- Primary (unique) index
- Secondary (nonunique) index
- Clustered index
- B*-tree index implementation
- Traversals of B*-tree index
- Examples

Index? What is it?

An **index** is a data structure that organizes data records on disk to optimize certain kinds of retrieval operations

An **index** is used to efficiently retrieve all records that satisfy a search condition on the search key fields of the index

An index is a function $f: K \rightarrow \wp(id_R)$ where K is set of keys and $\wp(id_R)$ is a powerset (a set of all sets) of identifiers (addresses) id_R of the records in a set R

Index? What is it?

Let *EMP* be a relational table over a relational schema *Employee(enumber, name, department)*

Then,

$F_{department}: \text{domain}(\text{department}) \rightarrow \wp(\text{id}_{EMP})$ is a function that maps the names of departments in $\text{domain}(\text{department})$ into the sets of identifiers of rows $\wp(\text{id}_{EMP})$ in relational table *EMP*

$F_{department}(d)$ returns the identifiers of all rows where a value of attribute *department* is equal to *d*

Introduction to Indexing

Index versus indexed file
organization



Index versus indexed file organization

- An **indexed file organization (index organized file)** is a function $f: K \rightarrow \wp(R)$ where K is a set of keys and $\wp(R)$ is a powerset (a set of sets) of records R
- An **index** maps a **value** into **set of row identifiers**
- An **index organized file** maps a **value** into a **set of records**
- A **relational table** can be **indexed** or it can be **index organized**
- An **indexed relational table** consists of several **index(es)** created separately from implementation of a **relational table** itself

Index versus indexed file organization

- An **index organized relational table** consists only of implementation of one **index** where an **index key** is the same as a **relational schema** of an **index organized table**
- Indexing in database systems is **transparent to data manipulation and data retrieval operations**
- It means that a database system automatically **modifies an index** and automatically decides whether an **index is used for search**

Introduction to Indexing

Primary (unique) index




Primary (unique) index

- A **primary (unique) index** is an index on a set of attributes equal to **primary** or **candidate key**
- A **primary index** is a function $f: K \rightarrow id_R$ where K is a set of **key values** and id_R is a set of identifiers (physical addresses) of rows in a relational table R
- **Primary index maps and index key into a single row identifier** (physical address of a row)
- A **primary key** in a **relational table** is always **automatically indexed** by a database system

Primary (unique) index

- For example, a relational table **EMPLOYEE** created over a relational schema *Employee(enum, name, department)* where **enum** is a **primary key** has an index automatically created on an attribute (**enum**)
- For example, a relational table **ENROLMENT** created over a relational schema *Enrolment(snumber, code, edate)* where (**snumber, code**) is a **primary key** has an index automatically created on a set of attributes (**snumber, code**)
- An index on (**snumber, code**) is a **composite index**

Primary (unique) index



$F_{\text{enum}}: \text{domain}(\text{enum}) \rightarrow \text{id}_{\text{EMPLOYEE}}$

Introduction to Indexing

Secondary (non-unique) index




Secondary (non-unique) index

- A **secondary index** is an index which is not **primary**
- For example, an index on an attribute (**name**) in a relational table **EMPLOYEE** created over a relational schema *Employee(enum, name, department)* is a **secondary (nonunique) index**
- For example, an index on a set of attributes (**name, department**) in a relational table **EMPLOYEE** created over a relational schema *Employee(enum, name, department)* is a **secondary index**

Secondary (non-unique) index

- For example, an index on an attribute (**snumber**) in a relational table **ENROLMENT** created over a relational schema **Enrolment(snumber, code, edate)** is a **secondary index**
- An index on a set of attributes (**enum, name**) in a relational table **EMPLOYEE** created over a relational schema ***Employee(enum, name, department)*** is still a **primary index** because (**enum, name**) is a **superkey**

Secondary (non-unique) index



$F_{\text{department}}: \text{domain}(\text{department}) \rightarrow \wp(\text{id}_{\text{EMPLOYEE}})$

Introduction to Indexing

Clustered index




Clustered index

- A **clustered index** is an index organized such that the ordering of rows is the same as ordering of keys in the index
- A clustered index is a function $f: K \rightarrow id_R$ where K is a set of keys and id_R is a set of row identifiers (addresses) in a relational table R such that $f(v)$ returns row identifier (address) of the first row in a sequence of rows such that a value of attribute K is equal to v
- Every **primary index** is **clustered**
- **Clustered index** provides faster access to data than **nonclustered secondary** index

Clustered index

- **Clustered index** has a very negative impact on performance of **INSERT** and **UPDATE** SQL statements
- Therefore, **clustered indexing** should be applied to mainly to **read-only data**

Clustered index



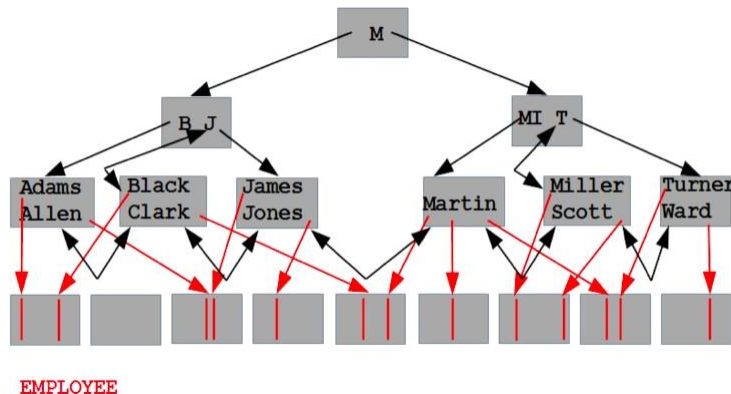
$f_{\text{department}}: \text{domain}(\text{department}) \rightarrow \text{id}_{\text{EMPLOYEE}}$

Introduction to Indexing

B*-tree index implementation



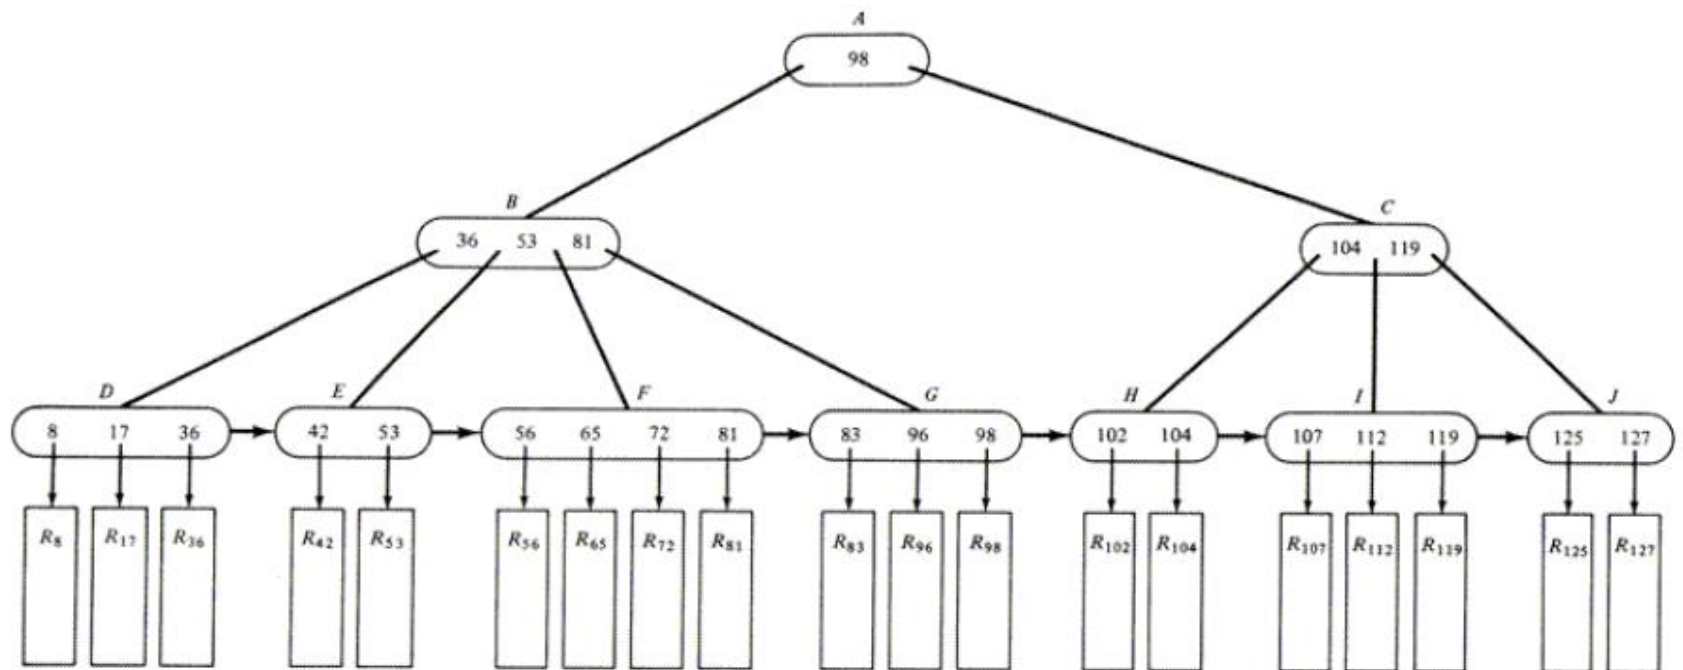
B*-tree index implementation



B*-tree can be traversed either:

- **vertically** from **root** to **leaf level** of a tree
- **horizontally** either from **left corner** of **leaf level** to **right corner** of **leaf level** or the opposite
- **vertically** and later on **horizontally** either towards **left lower corner** or **right lower corner** of **leaf level**

B*-TREE



Introduction to Indexing

Traversals of B*-tree index



Traversals of B*-tree index

- An index on a primary key (**enum**) in a relational table **EMPLOYEE** created over a relational schema *Employee(enum, name, department, salary)* is always built automatically by a database system
- A name of an index is the same as a name of primary key constraint in a relational table **EMPLOYEE**
- The following queries are processed through a **vertical traversal** of an index on (**enum**)

Traversals of B*-tree index

```
SELECT  *  
FROM    EMPLOYEE  
WHERE   enum=007;
```

```
SELECT  *  
FROM    EMPLOYEE  
WHERE   enum=007 AND department='mi6';
```

```
SELECT  enum  
FROM    EMPLOYEE  
WHERE   enum=007;
```

Traversals of B*-tree index

The following queries are processed through a **horizontal traversal** of leaf level of an index on **(enum)**

```
SELECT  count(*)  
FROM    EMPLOYEE;
```

```
SELECT  count(enum)  
FROM    EMPLOYEE;
```

```
SELECT  count(enum) /* Only if name is NOT NULL */  
FROM    EMPLOYEE;
```

Traversals of B*-tree index

The following queries are processed through a **horizontal traversal** of leaf level of an index on **(enum)**

```
SELECT  enum  
FROM    EMPLOYEE;
```

```
SELECT  enum, count(*)  
FROM    EMPLOYEE  
GROUP BY enum;
```

```
SELECT  enum  
FROM    EMPLOYEE  
ORDER BY enum;
```

Traversals of B*-tree index

Assume that we created an index on attribute **(name)** in a relational table **EMPLOYEE** created over a relational schema

Employee(enum, name, department, salary)

The following queries will be processed through a **vertical traversal** of an index on **(name)**

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     name='James';
```

Traversals of B*-tree index

```
SELECT  *  
FROM    EMPLOYEE  
WHERE   name = 'James' AND department = 'MI6';
```

```
SELECT  count(*)  
FROM    EMPLOYEE  
WHERE   name = 'James';
```

Traversals of B*-tree index

Assume that we created an index on attribute (**name, department**) in a relational table **EMPLOYEE** created over a relational schema

Employee(enum, name, department, salary)

The following queries will be processed through a **vertical traversal** of an index on (**name, department**)

```
SELECT      *  
FROM        EMPLOYEE  
WHERE       name='James' AND department='MI6';
```

Traversals of B*-tree index

```
SELECT    count(*)  
FROM      EMPLOYEE  
WHERE     name='James'  
AND       department='M16';
```

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     name='James'  
AND       department='M16'  
AND       salary>1000;
```


Traversals of B*-tree index

The following queries **can be processed** through a **vertical traversal** and later on **horizontal traversal** of an index on **(snum)**.

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     snum>300;
```

```
SELECT    count(*)  
FROM      EMPLOYEE  
WHERE     snum<007;
```

Traversals of B*-tree index

The following queries **can be processed** through a **vertical traversal** and later on **horizontal traversal** of an index on **(snum)**.

```
SELECT      *  
FROM        EMPLOYEE  
WHERE       snum>300 AND salary>1000;
```

Traversals of B*-tree index

Assume that we created an index on attribute (**name, department**) in a relational table **EMPLOYEE** created over a relational schema

Employee(enum, name, department, salary)

The following queries can be processed through a **vertical traversal** and later on **horizontal traversal** of an index on (**name, department**)

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     name>'James';
```

Traversals of B*-tree index

```
SELECT    count(*)  
FROM      EMPLOYEE  
WHERE     name<='James';
```

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     name='James' AND department>'MI6';
```

Traversals of B*-tree index

Assume that we created an index on attribute (**name, department**) in a relational table **EMPLOYEE** created over a relational schema

Employee(enum, name, department, salary)

The following queries can be processed through a **vertical traversal** and later on **horizontal traversal** of an index on (**name, department**)

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     name>'James' AND salary>1000;
```

Traversals of B*-tree index

```
SELECT  name, count(*)  
FROM    EMPLOYEE  
WHERE   name>'James' AND salary>1000  
GROUP BY name;
```

```
SELECT  *  
FROM    EMPLOYEE  
WHERE   name='James' AND salary>1000  
ORDER BY name;
```

Traversals of B*-tree index

Assume that we created an index on attribute (**name, department**) in a relational table **EMPLOYEE** created over a relational schema

Employee(enum, name, department, salary)

The following queries **can be processed** through a **horizontal traversal** of an index on (**name, department**)

```
SELECT    *  
FROM      EMPLOYEE  
WHERE     department='MI6';
```

Traversals of B*-tree index

```
SELECT *  
FROM   EMPLOYEE  
WHERE  department > 'MI6';
```

```
SELECT name, department  
FROM   EMPLOYEE;
```

```
SELECT name, department, count(*) F  
FROM   EMPLOYEE  
GROUP BY name, department;
```


Introduction to Indexing

Sample indexing



Sample indexing

What index should be created on a relational table
relational table **DEPARTMENT** created over a relational
schema

DEPARTMENT(dname, chairperson, budget) to
speed up the following queries?

```
SELECT    *  
FROM      DEPARTMENT  
WHERE     dname='MI6';
```

There is no need for any new index because an attribute
dname is a **primary key** and it is automatically indexed.

Sample indexing

```
SELECT *  
FROM DEPARTMENT  
WHERE dname='MI6' AND budget>10000;
```

There is no need for any new index because an attribute **dname** is a **primary key** and it is automatically indexed.

Sample indexing

What index should be created on a relational table **DEPARTMENT** created over a relational schema *DEPARTMENT(dname, chairperson, budget)* to speed up the following queries?

```
SELECT  *  
FROM    DEPARTMENT  
WHERE   budget=10000;
```

```
CREATE INDEX DEPT_IDX_BUDGET ON  
DEPARTMENT(budget);
```

Sample indexing

```
SELECT      *  
FROM        DEPARTMENT  
WHERE       budget=10000 AND chairperson='James';
```

```
CREATE INDEX DEPT_IDX_BC ON DEPARTMENT(budget,  
chairperson);
```

Sample indexing

```
SELECT      DISTINCT chairperson  
FROM        DEPARTMENT;
```

```
CREATE INDEX DEPT_IDX_CHAIR ON  
DEPARTMENT(chairperson);
```

Sample indexing

```
SELECT      *  
FROM        DEPARTMENT  
ORDER BY budget;
```

```
CREATE INDEX DEPT_IDX_BUDGET ON  
DEPARTMENT(budget);
```

Sample indexing

```
SELECT      chairperson, budget, count(*)  
FROM        DEPARTMENT  
GROUP BY    budget, chairperson;
```

```
CREATE INDEX DEPT_IDX_BC ON DEPARTMENT(budget,  
chairperson);
```


Sample indexing

```
SELECT      chairperson, budget, count(*)  
FROM        DEPARTMENT  
GROUP BY    chairperson, budget;
```

```
CREATE INDEX DEPT_IDX_CB ON  
DEPARTMENT(chairperson, budget);
```

References

- Elmasri R. and Navathe S. B., Fundamentals of Database Systems, Chapter 17 Indexing Structures for Files and Physical Database Design, 7th ed., The Person Education Ltd, 2017